

Atharva Tipre

Seattle, WA | 214.604.5921 | atharvatipre13@gmail.com | linkedin.com/in/atharva-tipre

SKILLS

Programming & Data Engineering: Python, SQL, PySpark, Spark SQL, R, Linux, ETL/ELT, Data Pipelines, Data Transformation, Data Validation, Data Quality, Data Profiling

Data Platforms & Warehousing: Snowflake, Amazon Redshift, MySQL, Dimensional Modeling, Star Schema, Data Marts, Data Lakehouse, CDC, Batch Processing, Real-Time Data Processing

Cloud & Infrastructure: AWS (Athena, S3, Glue, Lambda, EMR, Step Functions, SQS, SNS), Apache Spark, Apache Kafka, Git, JIRA, BitBucket

BI & Visualization: Tableau, Power BI, Alteryx, Adobe Analytics, Apache Airflow (MWAA), Dashboarding, KPI Reporting, Data Visualization, Executive Reporting

WORK EXPERIENCE

Sr. Analytics Engineer - AIML Solutions J.P. Morgan Chase *October 2024 – Present* Seattle, WA

- Designed and deployed a proactive data quality and observability framework across AWS Athena, Snowflake, and S3 — engineering **20+ production alerts** (volume change, zero traffic, schema drift) with automated threshold logic, reducing critical pipeline incidents by **50%** and JIRA turnaround from 10 to under 5 days across 4 AIML teams.
- Built a centralized SQL query repository with **10+** high-frequency queries embedded with source metadata, schema context, and sample outputs — reducing recurring data questions by **60–70%** and accelerating business reporting by **25%**.
- Led end-to-end Tableau Cloud migration for **100%** of AIML dashboards, resolving upstream connection reliability issues and cutting dashboard load times from **40–50s to under 10s**; migration playbook adopted as the org-wide standard by peer teams.
- Engineered **7+** data health monitoring dashboards achieving **>70%** pipeline coverage for AIML modeling teams, enabling sub-5-minute visibility into model performance, data freshness, and anomaly detection.
- Unified fragmented analytics pipelines across Athena, Snowflake, and S3 using Copilot AI — standardizing data contracts, KPI definitions, and reusable query templates enabling engineer self-service across Digital, Signals, UES, and Cross-Channel teams.

Analytics Engineer – Real Time Analytics J.P. Morgan Chase *July 2022 – September 2024* Seattle, WA

- Rebuilt cross-channel customer journey analytics by auditing raw data feeds, validating and reconciling data from disparate systems (clickstream, Adobe Analytics, channel events), and redesigning the full KPI framework — expanding from 5 to 10+ channels and growing analytics consumers from 2 to 15 business partners (~\$10M value).
- Built and deployed a propensity scoring model on 10M+ customer records using logistic regression (**80% accuracy**) — cutting campaign costs by **40%** and nearly doubling sign-up rates from 20–25% to ~50%; documented model features, labels, and schema definitions for downstream ML operationalization.
- Optimized Alteryx ETL transformation pipelines by adding a staging layer eliminating cascading timeouts, memory overruns, and retry storm failure modes — improving speed-to-market by **50%** across 20 downstream pipelines with ~\$100M revenue impact.
- Implemented automated data quality validation and anomaly detection controls monitoring completeness, accuracy, and consistency across large-scale customer datasets, reducing production reporting defects by **60%**.

Associate Analyst - Strategic Pricing Essilor Luxottica *July 2021 – June 2022* Dallas, TX

- Engineered an AWS S3 batch ETL pipeline and self-serve rebate investigation dashboard — cutting per-ticket investigation time by **>50%** and recovering ~\$120K in missed rebates within 4 months; scaled to 30K+ customer accounts across 3+ business groups.
- Developed customer segmentation and pricing analytics models using SQL and Python generating ~\$1.5M in projected promotional revenue and growing customer base by **20%**; prototyped Scorecards (analytical billing) for **30,000+** customers covering monthly product utilization and quarterly rebates.
- Designed interactive Tableau dashboards and KPI reports for executive and operational stakeholders surfacing pricing performance, utilization trends, and rebate analytics across 3 major business groups.

Data Analyst Incture L.L.C *May 2020 – August 2020* Pittsburgh, PA

- Built OCR-based Python/OpenCV data extraction pipelines transforming digitized and non-digitized documents into structured, analytics-ready datasets — reducing client operational costs by **30%**; presented findings on automated Tableau dashboards.

PROJECTS EXPERIENCE

- Data Engineering – Calendar Optimization Pipeline (AWS)**
Architected an end-to-end pipeline ingesting 1–100GB Calendly records into Amazon S3, processing with PySpark on Amazon EMR with Delta Lake for ACID-compliant storage, and orchestrating via Apache Airflow (MWAA) — achieving **50%+** reduction in processing time through optimized partitioning and cluster tuning.
- Applied Machine Learning – Bank Marketing Analysis (Banco de Portugal)**
Built and optimized 5 classification models on 50K customer records using scikit-learn with GridSearchCV hyperparameter tuning, achieving **91%** test accuracy; documented model logic, feature definitions, and usage examples for cross-functional stakeholders.
- Data Science – Instacart Market Basket Analysis**
Analyzed 3M+ grocery orders from 200K users using EDA and ensemble methods (LightGBM, XGBoost) to predict product reorder probability at **90%** accuracy, informing customer segmentation and targeting strategies.

EDUCATION

THE UNIVERSITY OF TEXAS AT DALLAS (M.S. Business Analytics - Data Science)

Dallas, TX

Dean's Excellence Scholarship, Naveen Jindal School of Management

MUMBAI UNIVERSITY (B.E. Electronics - Computer Engineering)

Mumbai, India